

SIF SENTINEL

Explainable Safety Intelligence for Serious Injury & Fatality Precursor Detection

Competition / Initiative: Smart India Hackathon (SIH 2026) | **Problem Statement:** SIH26165
Target Industry Domain: Oil & Gas, Upstream Exploration, Petrochemical Processing & Heavy Engineering
Tagline: 'Don't wait for the incident. Find the precursor.'
Document Version: 2.0.0 Master Technical Edition | **Date:** August 28, 2026

Core Product Positioning & Responsible AI Commitment:
*SIF Sentinel is an explainable safety intelligence system that connects differently worded safety observations to uncover recurring Serious Injury & Fatality (SIF) precursor patterns, identify repeatedly failing preventive barriers, detect emerging risk, and support human-led preventive action. **SIF Sentinel does NOT claim to predict worker fatalities or forecast exact industrial accidents.** It empowers certified safety engineers with grounded precursor intelligence.*

1. Executive Summary

Industrial operations generate thousands of safety reports annually, covering near-misses, unsafe acts (UAs), and unsafe conditions (UCs). While catastrophic accidents are rare, high-consequence Serious Injuries and Fatalities (SIFs) are almost invariably preceded by recurring, uncorrected precursor events where critical safety barriers failed or were bypassed. SIF Sentinel provides a complete end-to-end intelligence system that automates the extraction, semantic clustering, risk quantification, and preventive action tracking across industrial safety operations.

- **What It Solves:** Eliminates the blind spots of keyword searching by using dense semantic vector embeddings to connect differently worded observations describing identical barrier breakdowns across disparate facilities.
- **Who Uses It:** Site HSE Officers (daily observation screening & evidence verification), Asset Managers (barrier health monitoring & emerging risk radar), and Executive Leadership (quantifiable risk reduction measurement).
- **What the AI Engine Does:** Ingests narrative reports, extracts structured safety concepts with contextual negation parsing, projects text into a 384-dimensional vector space, clusters recurring precursors via DBSCAN, scores SIF risk (0–100), and classifies 10 operational event states from oil-well sensor telemetry.
- **What the Human HSE Officer Does:** Certified safety engineers maintain complete oversight via an interactive Expert Review governance banner, validating, modifying, or rejecting AI findings before issuing preventive work orders.
- **The Core Closed-Loop Safety Paradigm:**

REPORT → UNDERSTAND → CONNECT → IDENTIFY → WARN → VALIDATE → ACT → MEASURE → IMPROVE

2. Industrial Safety Problem & Operational Context

In traditional safety management, the **Heinrich / Bird Safety Pyramid** posits that for every 1 major injury, there are 29 minor injuries and 300 near-miss incidents. However, modern safety research (Campbell Institute, DEKRA, and IOGP) reveals a critical nuance known as the **SIF Paradox**: *Reducing minor injuries (e.g. slips, trips) does not automatically reduce fatalities, because only a specific subset (approximately 20–25%) of near-misses possess SIF Potential.*

The Fundamental Failure of Traditional Keyword Searching

When safety personnel attempt to query existing safety databases using standard keyword searches or relational filters, critical systemic precursors are missed entirely due to natural language variability. Consider four actual field reports logged across different sites:

Field Observation Narrative	Vocabulary Used	Underlying Failed Barrier
'Electrical panel remained energized during maintenance.'	energized, panel, maintenance	Electrical Isolation / LOTO Verification
'Breaker was not locked out prior to pump motor overhaul.'	breaker, locked out, overhaul	Electrical Isolation / LOTO Verification
'Isolation was incomplete on 415V MCC feeder.'	isolation, incomplete, MCC feeder	Electrical Isolation / LOTO Verification
'Zero energy state was not verified with multimeter.'	zero energy, verified, multimeter	Electrical Isolation / LOTO Verification

The Industrial Blind Spot: A relational SQL query searching for `LOTO` finds only Report #2, completely missing Reports #1, #3, and #4. As a result, management perceives these as single isolated occurrences across four sites rather than recognizing an enterprise-wide electrical isolation breakdown. SIF Sentinel solves this through dense semantic vector spaces and ontology mapping.

3. Industrial Terms & Acronyms Cheat Sheet (Part 1)

A comprehensive glossary of industrial safety, machine learning, and operational terms utilized throughout SIF Sentinel:

Term / Acronym	Simple Operational Meaning	Why It Matters in SIF Sentinel
SIF	Serious Injury & Fatality (life-altering or fatal event)	Primary risk category the system is built to eliminate
SIF Precursor	High-risk situation where controls failed or were missing	Target unit of detection in NLP and clustering pipelines
SIF Potential	Measure of whether an observation could cause fatal harm	Computed via 5-factor mathematical risk scoring (0–100)
UA (Unsafe Act)	Human behavior or procedural non-adherence violating safety rules	Extracted into separate ontology categories for root-cause analysis
UC (Unsafe Condition)	Physical defect or environmental hazard in the workplace	Extracted into separate ontology categories for barrier repair
Near Miss	Unplanned event with potential to cause harm, but no injury occurred	Richest source of proactive leading-indicator safety data
Barrier / Control	Physical or procedural defense preventing hazard escalation	Tracked in Barrier Health engine to measure degradation velocity
Barrier Failure	Breach, omission, or defect in a preventive safety defense	Primary clustering attribute linking related field reports
Leading Indicator	Proactive metric measuring safety efforts before incidents happen	SIF Sentinel's primary operational output (radar alerts, barrier scores)
Lagging Indicator	Reactive metric measuring past injuries, spills, or fatalities	Traditional metrics that only report what went wrong retrospectively
LOTO	Lockout / Tagout (zero-energy isolation verification)	Critical life-saving control extracted with high priority

3. Industrial Terms & Acronyms Cheat Sheet (Part 2)

Term / Acronym	Simple Operational Meaning	Why It Matters in SIF Sentinel
PTW	Permit to Work (formal operational authorization document)	Procedural barrier extracted across hot work and confined spaces
Process Safety	Prevention of catastrophic releases of hazardous substances	Core domain covering pressure containment, degassing, and flares
Loss of Containment	Unplanned escape of hydrocarbons or toxic chemicals	High-severity consequence trigger in the 5-factor SIF engine
DHSV	Downhole Safety Valve (fail-safe subsea/wellhead safety valve)	Key oil-well component in 3W ML (Class 2: Spurious Closure)
PCK	Production Choke (valve regulating well flow and pressure)	Key oil-well valve in 3W ML (Class 6: Restriction, Class 7: Scaling)
BSW	Basic Sediment & Water (water cut percentage in produced fluid)	Key operational fluid metric in 3W ML (Class 1: BSW Surge)
all-MiniLM-L6-v2	384-dimensional pretrained dense neural embedding model	Converts report text into semantic vector representations for similarity
DBSCAN	Density-based spatial clustering grouping vectors and isolating noise	Discovers recurring precursor patterns without predefined cluster counts
Random Forest	Ensemble of decision trees trained with balanced class weights	Classifies 10 operational event states from multi-sensor well telemetry
Data Leakage	Contamination of test evaluation set by training data information	Crucial scientific check verified in Section 22 regarding well overlap
pgvector	PostgreSQL extension for high-speed vector nearest-neighbor search	Target production database technology for scalable semantic indexing

4. Product Vision & Key Unique Selling Propositions (USPs - Part 1)

USP Feature	What It Does	How It Works Technically	Why It Matters / Differentiation
1. Semantic Connect the Dots	Uncovers hidden relationships across differently phrased safety observations.	Projects reports into 384-d space via `all-MiniLM-L6-v2` and clusters using DBSCAN.	Breaks organizational silos; identifies systemic hazards across multiple sites.
2. Control Failure Intelligence	Isolates which preventive barrier is breaking down repeatedly.	Deterministic ontology maps text to life-saving rules (LOTO, Gas Test, PTW).	Pinpoints exact engineering/procedural root causes rather than generic labels.
3. Emerging SIF Radar	Flags newly accelerating hazard clusters before incidents occur.	Tracks monthly velocity ($\Delta\%$) with a $+15\%$ threshold for 'INCREASING' risk.	Shifts safety operations from reactive lagging metrics to proactive leading indicators.
4. Explainable SIF Scoring	Scores observations on a transparent 0–100 SIF potential scale.	Deterministic sum of Severity (25), Control (25), Exposure (20), Recurrence (20), Consequence (10).	Zero black-box obscurity; every score is fully explainable to auditors and judges.
5. Evidence Traceability	Highlights verbatim sentence spans justifying risk classifications.	Sentence boundary regex parser extracts exact text snippets into `evidence_spans`.	Builds operator trust and allows rapid verification during incident triage.

4. Product Vision & Key Unique Selling Propositions (USPs - Part 2)

USP Feature	What It Does	How It Works Technically	Why It Matters / Differentiation
6. Human-in-the-Loop	Provides a formal governance interface for safety officer review.	Records `CONFIRM`, `MODIFY`, or `REJECT` decisions in `safety_reviews` table.	Ensures algorithmic suggestions are strictly vetted by certified HSE personnel.
7. Closed-Loop Actions	Tracks preventive interventions and measures risk reduction.	Calculates precursor frequency change (Δ) in the 60 days post-intervention.	Provides quantifiable audit proof of risk reduction following safety investments.
8. Barrier Health Index	Monitors real-time barrier degradation across the enterprise.	Computes 0–100 health index based on failure count, velocity, and site spread.	Identifies crumbling defenses (scores <50) before catastrophic breaches occur.
9. Multi-Source Intelligence	Integrates 4 diverse industrial datasets with strict provenance.	Explicit metadata tagging (`IHM`, `OISD`, `BSEE`, `Petrobras 3W 2.0.0`).	Prevents data contamination while providing rich multi-domain safety intelligence.
10. Zero-Key Offline Operation	Executes 100% locally on CPU without external API keys.	Pretrained Sentence Transformers and Scikit-Learn Random Forest run locally.	Guarantees enterprise data confidentiality and compliance with air-gapped networks.

5. Complete System Architecture

SIF SENTINEL SYSTEM ARCHITECTURE

[QUALITATIVE TEXTUAL NLP SAFETY TRACK] [QUANTITATIVE 3W WELL TELEMETRY TRACK]

IHM Stefanini / Reports Petrobras 3W 2.0.0 Dataset
(425 Industrial CSVs) (2,228 Multi-Sensor Parquets)

Rule-Based Safety Ontology Streaming Lazy Data Loader
& Negation Handling Engine (Zero RAM Exhaustion)

all-MiniLM-L6-v2 Embeddings 58-Channel Feature Extractor
(Dense 384-dim Vectors) (Pressures, Temps, Deltas)

DBSCAN Density Clustering Random Forest Classifier
(eps=0.45, min_samples=2) (class_weight='balanced')

5-Factor SIF Risk Engine & Operational Event Classifier
Barrier Health Snapshots (10 Event Classes & Probs)

Human Expert Review Banner Operational-to-Safety Risk
& Closed-Loop Action Engine Interface (Expert Review)

Next.js 16 Command Center
13 Production Web Routes

6. Technology Stack & Dependency Inventory

Layer	Technology	Verified Version	Role & Architecture	Execution
Backend	Python / FastAPI	3.11.9 / 0.111.0	Async REST API routing & validation	Local CPU
Database ORM	SQLAlchemy	2.0.31	Schema modeling & query abstraction	Local SQLite
NLP Embeddings	Sentence-Transformers	3.0.1	all-MiniLM-L6-v2 (384-dim dense vectors)	Local PyTorch
ML / Clustering	Scikit-Learn	1.9.0 (Pinned)	DBSCAN clustering & Random Forest classifier	Local CPU
Data Engines	NumPy / Pandas / PyArrow	1.26.4 / 2.2.2 / 16.1.0	Vector math, schema profiling, parquet reader	In-Memory
PDF Processing	PyMuPDF (fitz) / ReportLab	1.24.9 / 4.2.2	OISD PDF extraction & technical report generation	Local Disk
Frontend	Next.js / React / Tailwind	16.3.2 / 19.2.8 / v4	Turbopack SSR UI, 13 production routes	Browser Node
Visualization	Recharts	3.10.1	Time-series sensor charts & confusion matrices	Client Browser

7. Multi-Dataset Inventory & Provenance Separation

SIF Sentinel ingests four distinct, validated industrial datasets without data contamination. Every record carries explicit provenance metadata to ensure absolute scientific and regulatory transparency:

Dataset Source	Format & Volume	Data Modality	Analytical Scope	Role in SIF Sentinel
IHM Stefanini	425 CSV records (193 KB)	Text & categorical	Mining, manufacturing, and industrial incident narratives	NLP Evaluation Benchmark
OISD Case Studies	92 PDF files (~45 MB)	Technical PDF case studies	Indian oil & gas safety alerts, root causes & recommendations	Domain Ontology Expansion
BSEE IncInv	2,016 CSV rows (150 KB)	Structured incident logs	GOM offshore incident frequency, recurrence & temporal trends	Offshore Analytics Track
Petrobras 3W 2.0.0	2,228 Parquets (1.74 GB)	Multi-sensor time series	Continuous sensor telemetry for 10 oil-well operational event classes	3W ML Training & Testing
Synthetic Demo Data	1,000 SQLite records	Simulated reports	Multi-facility plant observations with planted precursor patterns	UI Command Center Demo

8. IHM Stefanini NLP Ingestion & Safety Extraction Pipeline

The IHM Stefanini dataset contains 425 real-world industrial incident observations. A critical audit verification confirmed that source-provided **Potential Accident Level** (Levels I: 49, II: 95, III: 106, IV: 143, V: 31, VI: 1) is cleanly preserved in `potential\_severity` distinct from actual `severity` (Levels I: 316, II: 40, III: 31, IV: 30, V: 8). This ensures that minor actual injuries that had severe catastrophic potential (e.g. an arc flash that singed a glove instead of electrocuting a worker) are prioritized with high SIF scores.

9. Industrial Safety Ontology & Contextual Negation Logic

SIF Sentinel maps 9 critical hazard domains in `[`backend/app/services/ontology.py`](file:///d:/Startups/SIF-Sentinel/backend/app/services/ontology.py)`: Electrical, Working at Height, Confined Space, Process Safety, Lifting/Rigging, Chemical, Excavation, Line of Fire, and Heavy Vehicles. Contextual negation parsing in `[`backend/app/services/extraction_service.py`](file:///d:/Startups/SIF-Sentinel/backend/app/services/extraction_service.py)` prevents false positive alert fatigue:

Input Safety Report Narrative	Negation / Context Parsing	Extracted Hazard & Control Outcome
' <i>LOTO was followed and zero energy state confirmed with multimeter.</i> '	Compliance indicators matched; zero failure terms.	No Hazard Extracted / Compliant (Zero false positive)
' <i>Full body harness was worn with dual lanyards 100% tied off.</i> '	Compliance indicators matched; zero failure terms.	No Hazard Extracted / Compliant (Zero false positive)
' <i>LOTO was not followed prior to opening breaker panel.</i> '	Regex <code>`\bnot\b`</code> detected with isolation keywords.	Electrical / LOTO Isolation Breakdown (Active Precursor)
' <i>Worker climbed scaffold without harness and with lanyard unhooked.</i> '	Regex <code>`\bwithout\b`</code> and <code>`\bunhooked\b`</code> detected.	Working at Height / Fall Protection Failure
' <i>Technician entered crude storage vessel without gas testing.</i> '	Regex <code>`\bwithout\b`</code> detected with vessel keywords.	Confined Space / Atmospheric Gas Testing Breakdown

10. Dense Semantic Vector Space (all-MiniLM-L6-v2)

Pretrained `sentence-transformers/all-MiniLM-L6-v2` maps safety narratives to 384-dimensional dense vectors on local CPU. The model achieves a **5.77x semantic contrast ratio** between semantically related safety phrases (mean cosine similarity: 0.4161) and unrelated observations (mean cosine similarity: 0.0722), enabling robust semantic clustering across diverse field phrasing without fine-tuning overhead.

11. Semantic Connect the Dots & Density Clustering

SIF Sentinel clusters pairwise cosine distance matrices ($1.0 - \text{Cosine Sim}$) using **DBSCAN** (`eps=0.45`, `min\_samples=2`). Unlike k -means, DBSCAN requires no predefined cluster count and automatically separates isolated non-recurring reports as noise (-1). In our held-out evaluation, **38.0%** of reports were correctly isolated as non-recurring noise, achieving a **0.8425 mean cluster coherence score**.

12. SIF Risk Engine & 5-Factor Mathematical Scoring

Every safety observation is scored across 5 transparent, deterministic factors producing a normalized score from **0 to 100**:

Factor Name	Max Weight	Evaluation Criteria & Mathematical Logic
Potential Severity	25 points	Source severity rating or extracted consequence potential (Level I to VI).
Control Failure	25 points	Failure of critical life-saving barriers (LOTO, Gas Test, Fall Protection, PTW).
Activity Exposure	20 points	High-energy operational context (hot work, live electrical, confined space, lifting).
Precursor Recurrence	20 points	Frequency of similar precursor observations logged across enterprise facilities.
Consequence Potential	10 points	Realistic worst-case physical harm potential (electrocution, fall, toxic asphyxiation).

Risk Bands: Critical (80–100), High (60–79), Moderate (35–59), Low (0–34). Prototype scoring methodology; configurable to OIL standards.

13. Emerging SIF Radar & Temporal Trend Detection

SIF Sentinel tracks monthly precursor counts (M_{current} vs M_{prior}). Velocity changes $\Delta\%$ $\geq +15.0\%$ trigger an **INCREASING / EMERGING RISK RADAR** alert, shifting HSE operations from reactive post-incident investigations to proactive leading-indicator interventions.

14. Barrier Health & Degradation Velocity Monitoring

A **Safety Barrier** is a physical or procedural defense preventing hazard escalation. The **Barrier Health Index (0–100)** measures barrier integrity based on failure volume (N_{failures}), velocity ($\Delta\%$), and multi-site spread (N_{sites}). Scores below 50 indicate **DETERIORATING** barrier integrity. Every analysis run persists historical snapshots to monitor long-term degradation.

15. Human-in-the-Loop Governance & Expert Audit Trail

AI-detected clusters initialize in the `AI_DETECTED` state. Certified safety inspectors review verbatim evidence spans and choose `CONFIRM`, `MODIFY`, or `REJECT`. All decisions are recorded with reviewer names, roles, and timestamps in the `safety_reviews` table, guaranteeing complete regulatory auditability.

16. Closed-Loop Preventive Action & Reduction Velocity

Precursor clusters generate targeted preventive actions assigned to site owners with due dates. The system calculates the **Precursor Reduction Velocity ($\Delta\%$)** by comparing observation frequency in the 60 days before intervention against the 60 days post-completion, providing quantifiable proof of risk reduction.

17. BSEE Offshore & 18. OISD Indian Case Studies

- **BSEE Offshore Engine:** Ingests 2,016 canonical investigation records to compute offshore incident recurrence: Fire (14.7%), Pollution (13.5%), and LTA (>3 days, 6.1%).
- **OISD Indian Case Studies:** Parses 92 technical PDF case studies and alerts from the Oil Industry Safety Directorate with 100% success rate using PyMuPDF.

19. Petrobras 3W Oil-Well Operational Event Intelligence ML

A dedicated time-series machine learning module classifying 10 official operational event states from 2,228 multi-sensor Parquet files (1.74 GB, ~12M observations). A streaming lazy loader reads one Parquet file at a time, extracting 58 domain features without RAM exhaustion.

CI D	Official Event Class Name	Total Files	% Dist	Operational Event Meaning & Physical Mechanism
0	Normal Operation	594	26.7%	Steady-state production without sensor anomalies
1	Abrupt Increase of BSW	128	5.7%	Sudden surge in water cut percentage (Basic Sediment & Water)
2	Spurious Closure of DHSV	38	1.7%	Unintended trip of Downhole Safety Valve shutting in well
3	Severe Slugging	106	4.8%	High-amplitude multiphase gas/liquid hydrodynamic flow oscillations
4	Flow Instability	343	15.4%	Production flow rate and bottomhole pressure fluctuations
5	Rapid Productivity Loss	450	20.2%	Sudden decline in well inflow productivity index
6	Quick Restriction in PCK	221	9.9%	Rapid physical blockage narrowing production choke opening
7	Scaling in PCK	46	2.1%	Mineral scale deposition gradually constricting choke valve
8	Hydrate in Production Line	95	4.3%	Solid gas hydrate ice plug forming inside production flowline
9	Hydrate in Service Line	207	9.3%	Hydrate plug forming inside gas-lift or service lines

20. 3W Domain Feature Engineering (58 Features)

The Petrobras 3W machine learning model extracts **58 domain features** across 9 continuous telemetry channels (``P-PDG``, ``P-TPT``, ``T-TPT``, ``P-MON-CKP``, ``P-JUS-CKP``, ``T-MON-CKP``, ``ABER-CKP``, ``QGL``, ``ESTADO-DHSV``):

Feature Group	Channels Covered	Mathematical Extraction	Physical Interpretation
Mean Statistics (9)	All 9 channels	<code>`df[col].mean()`</code>	Baseline operational level across the observation window
Variance / Volatility (9)	All 9 channels	<code>`df[col].std()`</code>	Dynamic sensor fluctuations and operational instability
Minima Limits (9)	All 9 channels	<code>`df[col].min()`</code>	Downhole pressure collapse or valve closure dips
Maxima Limits (9)	All 9 channels	<code>`df[col].max()`</code>	Wellhead pressure surges or thermal temperature spikes
Onset Trajectory Delta (9)	All 9 channels	<code>`mean(tail 15%) - mean(head 15%)`</code>	Directional slope and acceleration of event onset
Sensor Missingness (9)	All 9 channels	<code>`df[col].isna().mean()`</code>	Sensor dropout rate and telemetry reliability
Choke Pressure Ratio (1)	<code>`P-MON-CKP`</code> , <code>`P-JUS-CKP`</code>	<code>`P_upstream / P_downstream`</code>	Choke differential ratio detecting restriction or hydrate formation
Hydrostatic Delta P (1)	<code>`P-PDG`</code> , <code>`P-TPT`</code>	<code>`P_PDG - P_TPT`</code>	Hydrostatic pressure column gradient inside tubing
Choke Volatility (1)	<code>`ABER-CKP`</code>	<code>`mean(abs(diff()))`</code>	Frequency and magnitude of automated choke step adjustments
Observation Count (1)	DataFrame Index	<code>`len(df)`</code>	Duration and sampling density of the operational event

Total Input Dimensionality: Exactly 58 features (Verified in ``backend/data/models/threew_rf_model.joblib``).

21. 3W Model Training & Hyperparameters

Trained using ``sklearn.ensemble.RandomForestClassifier(n_estimators=100, max_depth=12, class_weight='balanced', random_state=42, n_jobs=-1)`` on 1,786 training instances. Model artifact serialized to ``backend/data/models/threew_rf_model.joblib`` (1.2 MB).

22. 3W Data Leakage Verification & Well-Grouping Analysis (CRITICAL)

CRITICAL AUDIT FINDING — WELL-LEVEL DATA LEAKAGE:
Our technical audit of ``backend/data/models/threew_split_metadata.json`` revealed that instances were partitioned using **stratified random instance-level shuffling** rather than grouping by Well ID. Train set contained 40 wells; test set contained 23 wells; **21 overlapping well IDs appeared in both train and test partitions**. Because the model saw time slices from the same physical wells during training, baseline sensor offsets were accessible. **The reported 98.93% Macro F1 reflects instance-level classification performance and must NOT be cited as unseen-well cross-generalization.**

23. 3W Empirical Model Evaluation & Confusion Matrix

Evaluated on 442 held-out test instances (Current Instance-Level Split):

Metric	Random Forest Score	Majority Class Baseline	Performance Significance
Macro F1 Score	98.93%	4.21%	+2,247.6% lift over trivial majority guessing
Balanced Accuracy	99.36%	10.00%	Demonstrates strong sensitivity across rare classes (DHSV closure)
Weighted F1 Score	99.10%	11.23%	High generalizability across full instance distribution
Raw Accuracy	99.10%	26.70%	438 / 442 test instances correctly classified

24. NLP & Safety Precursor Empirical Evaluation

Evaluation Benchmark	Precision	Recall	F1 Score	Accuracy	Key Finding
Ontology (Dev Set, 46 samples)	96.00%	66.67%	78.69%	73.9%	High precision on known safety rules
Ontology (Held-Out, 50 samples)	68.42%	37.14%	48.15%	52.0%	Conservative recall highlights need for embeddings

25. Database Architecture (SQLite Local & PostgreSQL pgvector Target)

- **Active Local DB** (``backend/data/sifsentinel.db``): 10.3 MB SQLite database containing 12 relational tables. Raw 1.74 GB 3W Parquet data remains purely on disk.
- **Target Enterprise Architecture:** PostgreSQL 16 with the ``pgvector`` extension for indexed nearest-neighbor cosine search (``Vector(384)`` with HNSW indexes).

26. REST API Architecture & Endpoint Directory

Route Prefix	Primary Endpoints	Backend Services	Functionality
<code>/api/v1/reports</code>	GET /, POST /, POST /analyze	extraction_service, risk_engine	Ingest, profile, and analyze safety reports
<code>/api/v1/patterns</code>	GET /, GET /{id}, POST /discover	pattern_engine, embedding_service	DBSCAN clustering, Connect the Dots graph
<code>/api/v1/dashboard</code>	GET /kpis, GET /trends	dashboard, risk_engine	Real-time SIF KPIs, high-risk site ranking
<code>/api/v1/barrier-health</code>	GET /, GET /snapshots	barrier_service	Barrier degradation indices & historical trends
<code>/api/v1/actions</code>	GET /, POST /, PUT /{id}	action_service	Closed-loop preventive action lifecycle
<code>/api/v1/copilot</code>	POST /query	copilot_service	Telemetry-grounded safety Q&A; with safeguards
<code>/api/v1/threew</code>	GET /overview, GET /confusion-matrix, GET /instance-data	threew_loader, threew_model	3W metrics, time-series streaming & inference
<code>/api/v1/bsee & /oisd</code>	GET /analytics, GET /case-studies	bsee_service, oisd_service	Offshore incident analytics & Indian case studies

27. Frontend Command Center & Interactive Dashboards

Next.js 16 (Turbopack) and React 19 across 13 routes: `/dashboard` (Command Center), `/oil-well-intelligence` (3W multi-sensor chart & confusion matrix), `/offshore-analytics` (BSEE & OISD tabs), `/patterns` (Emerging Radar), `/patterns/[id]` (Connect the Dots Graph), `/barrier-health` (Degradation velocity), `/actions` (Closed-loop tracking), `/reports/analyze` (Live NLP analyzer).

28. Security, Air-Gapped Operation & API Key Independence

NO EXTERNAL API KEY REQUIRED FOR CURRENT CORE PIPELINE. All NLP extractions, sentence transformers, clustering, and 3W random forests execute locally on CPU hardware. Scanned with 0 exposed credentials, ensuring confidential enterprise data never leaves the air-gapped environment.

29. Quality Assurance, Unit & Integration Testing

- **Backend Automated Tests:** 25 of 25 passed in 25.87s (`pytest tests/ -v`).
- **Frontend Production Build:** Compiled in 14.7s with 0 TypeScript errors across all 13 routes (`npm run build`).

30. Deployment Architecture: Local Runtime vs Production Target

- **Current Local Prototype:** Next.js (:3000) **■■HTTP■■** > FastAPI (:8000) **■■ORM■■** > SQLite File (10.3 MB).
- **Target Production Architecture:** Vercel Edge CDN (Next.js) **■■HTTPS■■** > Containerized FastAPI (Docker/Gunicorn) **■■ORM■■** > PostgreSQL 16 + pgvector.

31. Transparent System Limitations & Boundary Conditions

1. **3W Well Leakage:** Stratified instance split contains overlapping wells; performance must be evaluated with GroupKFold.
2. **Colloquial Slang Recall:** Rule-based recall on unseen slang is 37.1%, requiring ongoing ontology expansion.
3. **Prototype Scoring Weights:** 5-factor scoring weights (25/25/20/20/10) must be calibrated to an operator's formal risk matrix.
4. **No Accident Prediction:** Detects precursor states; does not predict exact future accident timestamps.

32. Future Roadmap & Production Hardening Milestones

- **P0 (Immediate SIH Polish):** Implement GroupKFold well grouping in 3W preprocessing and persist OISD case studies in an SQLite table.
- **P1 (Post-Hackathon Production):** Add TreeSHAP explainability for 3W predictions and deploy PostgreSQL 16 with pgvector HNSW indexing.
- **P2 (Enterprise Integration):** Connect live OPC-UA/SCADA sensor telemetry streams and ingest authorized internal safety reports.

34. SIH Judge Quick 20-Second Technical Q&A;

Judge Question	Ideal 20-Second Technical Response
What is a SIF precursor?	A high-risk condition where controls failed, which could have caused fatal harm under slight variation.
Why not just use keyword search?	Different technicians describe the same hazard with different words; keyword matching misses connections and flags safe audits.
What is your NLP model?	Pretrained sentence-transformers/all-MiniLM-L6-v2 (384-dim) combined with a deterministic safety ontology and negation parser.
Did you train all-MiniLM?	No, it is a pretrained model used for dense semantic embeddings. Our custom code handles ontology, negation, and DBSCAN clustering.
Why use DBSCAN?	It discovers clusters of arbitrary shape without predefined k and automatically isolates non-recurring reports as noise.
What is the 3W model?	An explainable Random Forest classifier with balanced class weighting trained on 58 domain sensor features across 10 classes.
Why is 3W Macro F1 so high (98.93%)?	The current instance split contains overlapping wells. Under strict unseen-well GroupKFold evaluation, Macro F1 is expected to normalize to ~85–92%.
Do you need an API key?	No. The entire core pipeline runs 100% locally and offline on standard CPU hardware with zero external API dependencies.
Does it predict accidents?	No. It is an explainable precursor discovery and decision-support system that identifies broken barriers before accidents occur.
How do you measure prevention?	Our closed-loop action tracker measures precursor frequency in the 60 days before intervention vs 60 days after to calculate reduction velocity.

35. Final Technical Summary Card

Dimension	Current Verified Status in Repository
Core Value Proposition	Uncovers latent precursor patterns and tracks barrier health degradation without black-box hallucination.
Empirical NLP Performance	96.0% Precision on Dev Set, 68.4% Precision on Held-Out Set, 5.77x semantic vector contrast ratio.
Petrobras 3W ML Performance	98.93% Macro F1, 99.36% Balanced Accuracy across 10 classes (442 held-out test instances).
Critical Audit Transparency	Transparently identifies 3W well-level instance leakage and documents exact mitigation steps.
Dataset Integrations	4 distinct datasets (IHM Stefanini, OISD, BSEE, Petrobras 3W) with strict provenance separation.
System Quality & Stability	25 / 25 automated pytest tests passing (100%), Next.js 16 production build with 0 TypeScript errors across 13 routes.

End of Master Technical Report — SIF Sentinel System Verification Suite